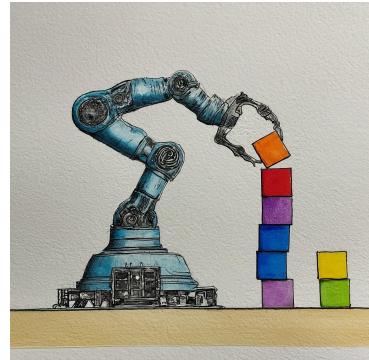


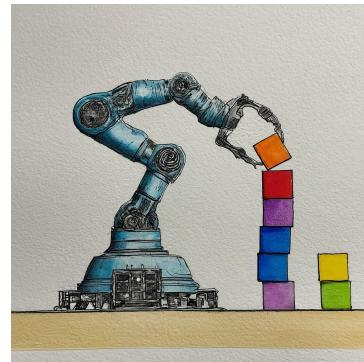
Learning for Integrated Task and Motion Planning

2025 AAAI Bridge Program



Content

- About the program
- Schedule overview
- An introduction to integrated task and motion planning

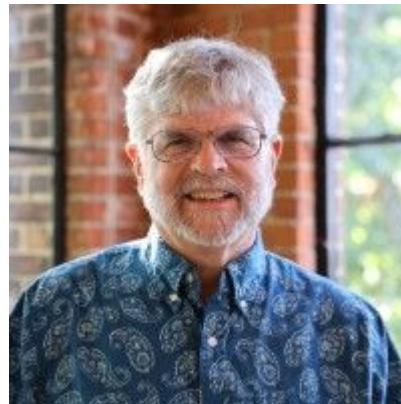


Organizing Team

Sarah Keren



Brian Williams



Michael Posa



Collaborative AI and Robotics



Model-based Embedded and
Robotic Systems



Dynamic Autonomy and Intelligent
Robotics Lab



Our github links



https://github.com/CLAIR-LAB-TECHNION/AAAI_25_Bridge_TMP

<https://github.com/CLAIR-LAB-TECHNION/CLAIR-TMP-Tutorials>

Schedule

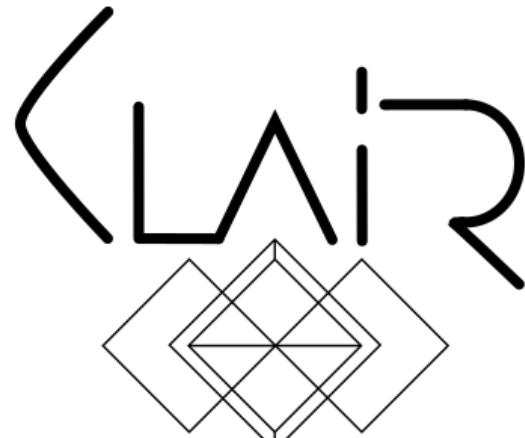
Day 1

Time	Session
9:00 AM - 10:30 AM	Intro + Tutorial 1 - Sarah Keren (Technion)
10:30 AM - 11:00 AM	Break
11:00 AM - 12:30 PM	Brian Williams (MIT)
12:30 PM - 2:00 PM	Lunch (on your own; no sponsored lunch provided)
2:00 PM - 3:30 PM	Jiayuan Mao (MIT) + Tutorial 2 -Sarah Keren (Technion)
3:30 PM - 4:00 PM	Break
4:00 PM - 5:30 PM	Michael Posa (UPenn) + Laser Talks

Day 2

Time	Session
9:00 AM - 10:30 AM	Tutorial 3 - Sarah Keren (Technion)
10:30 AM - 11:00 AM	Break
11:00 AM - 12:30 PM	David Held (CMU)
12:30 PM - 2:00 PM	Lunch (no sponsored lunch provided)
2:00 PM - 3:30 PM	Peter Stone (University of Texas at Austin and Sony AI)
3:30 PM - 4:00 PM	Break
4:00 PM - 5:30 PM	Panel: Peter Stone, David Held, Michael Posa, Brian Williams, Rachel Holladay (UPenn), Sarah Keren

A Very Brief Intro to Task and Motion Planning



Collaborative AI and Robotics

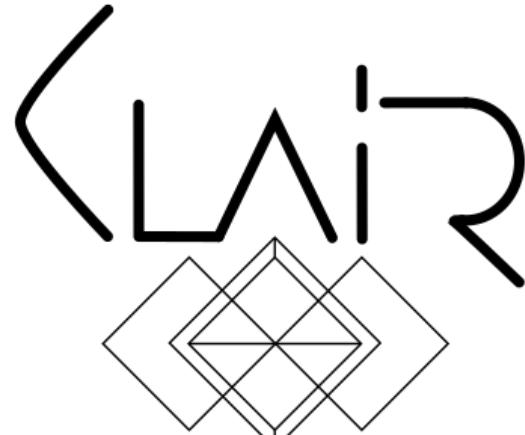
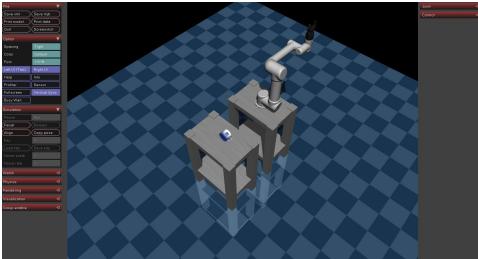
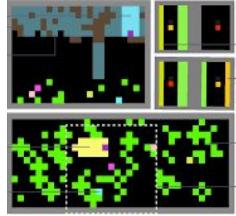
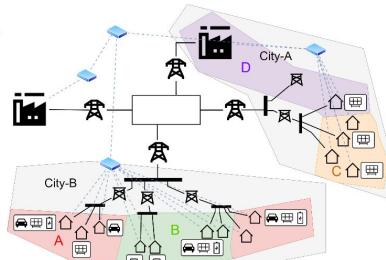
Sarah Keren



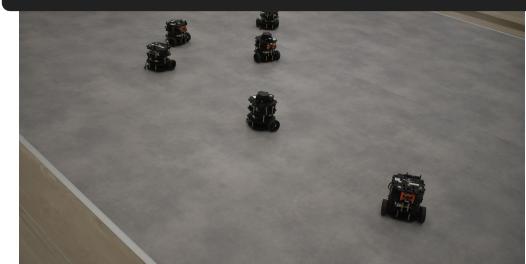
TECHNION



The Henry and Marilyn Taub
Faculty of Computer Science



Collaborative AI and Robotics



Sarah Keren



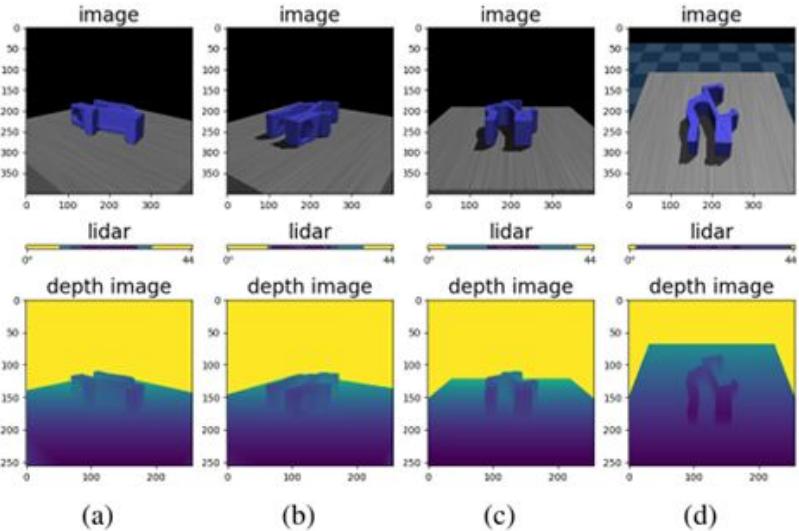
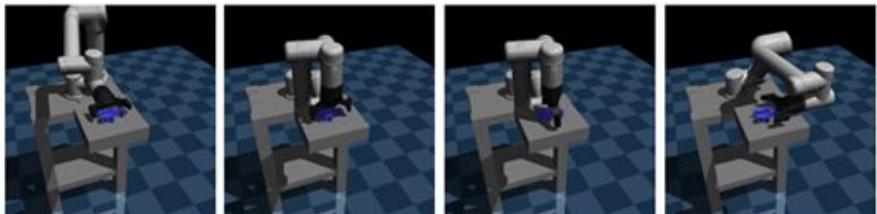
TECHNION |



The Henry and Marilyn Taub
Faculty of Computer Science

One of our projects

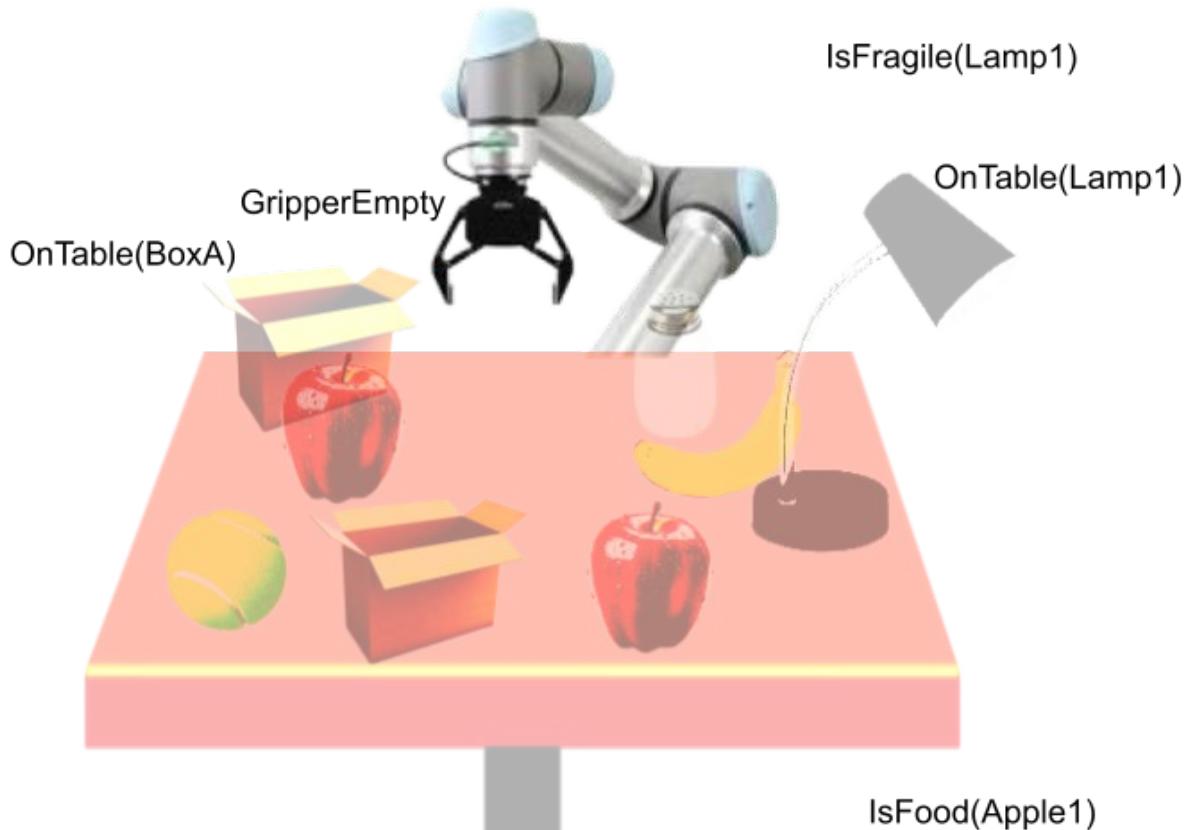
Value of Assistance for long term decision making under uncertainty

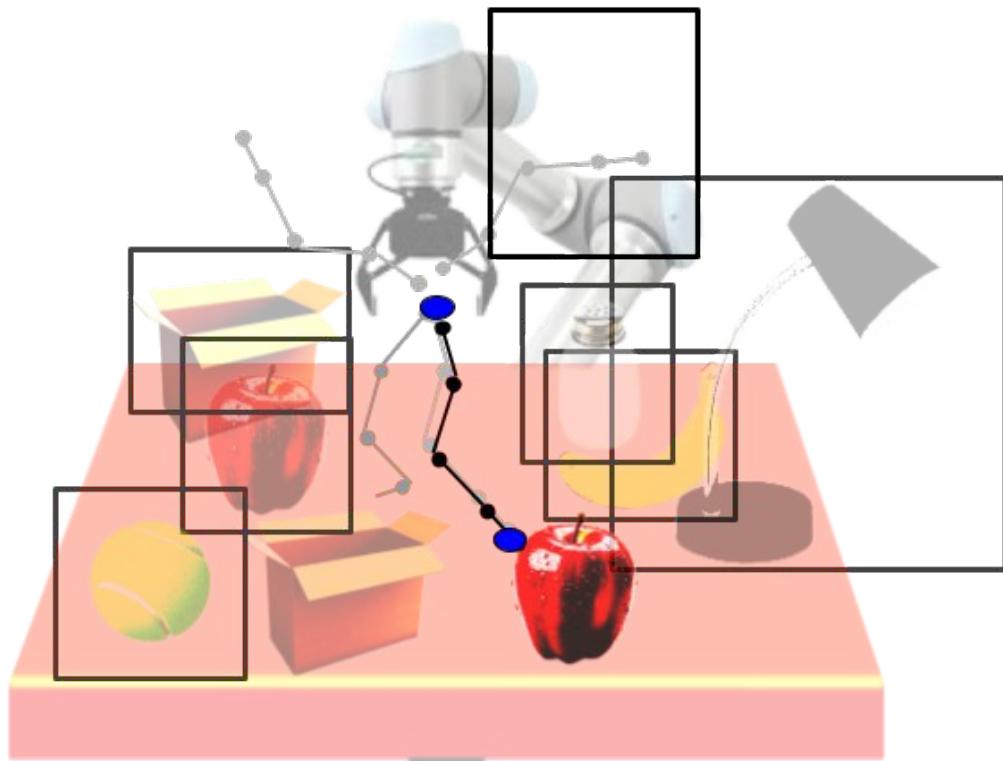


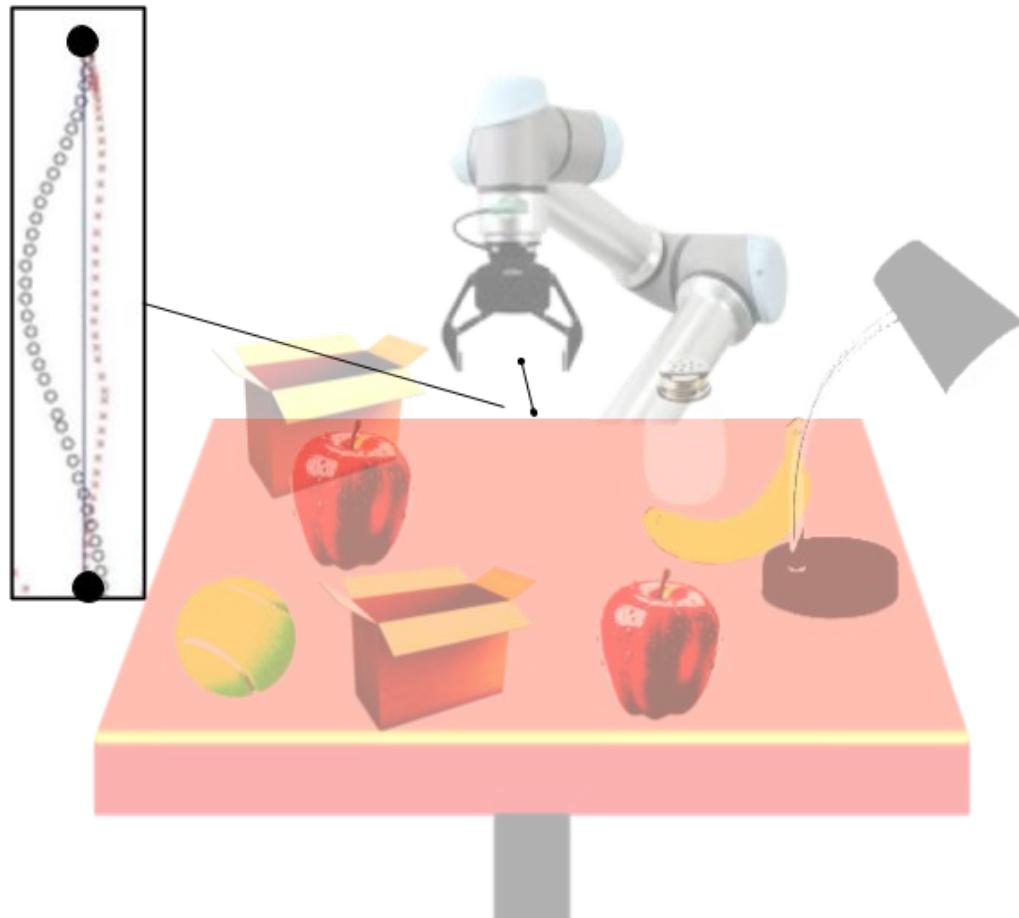
Example

“I am hungry”









Outline

Motivation for integrating task and motion planning (and control)

Components of a robotic agent

Task planning and its limitations

Motion planning and its limitations

Integrating task and motion planning

Let's (of course) start with a blocks world example

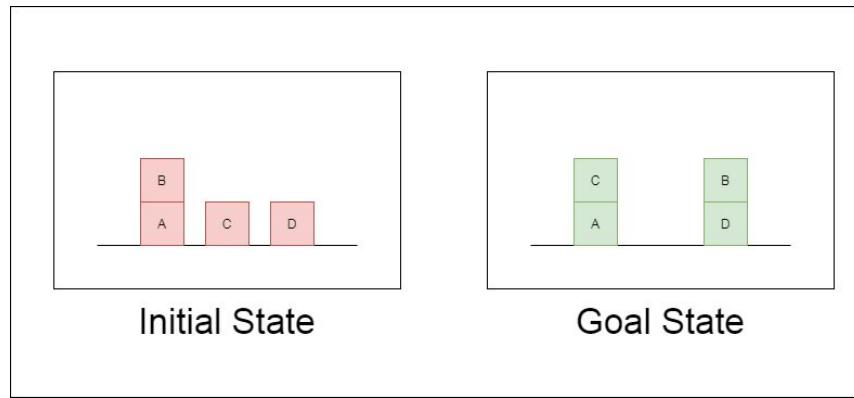
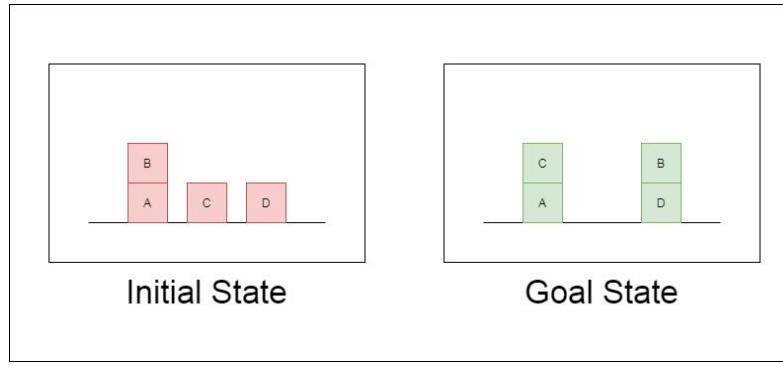


Image from <https://apoorvdixit619.medium.com/goal-stack-planning-for-blocks-world-problem-41779d090f29>

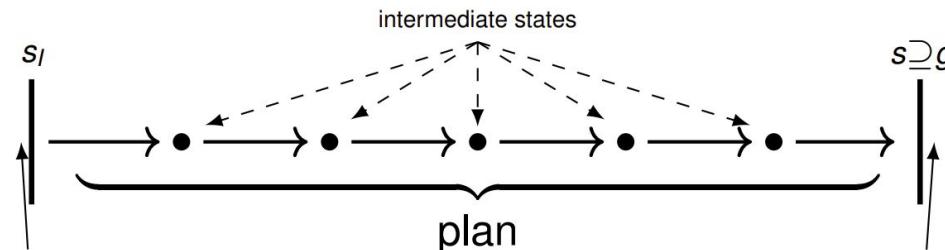


```
(:types block)
(:predicates (on ?b1 ?b2 - block)
             (clear ?b - block)
             (handempty)
             (holding ?b - block))

(:action stack
  :parameters (?b1 ?b2 - block)
  :precondition (and (holding ?b1) (clear ?b2))
  :effect
    (and (not (holding ?b1))
         (not (clear ?b2))
         (clear ?b1)
         (handempty)
         (on ?b1 ?b2)))
```

```
(:init
  (handempty) (ontable D) (clear D) (on B A) ...)

(:goal (and
  (on B D) (ontable D)
  (on C A) (ontable A))))
```



description of the
initial world situation

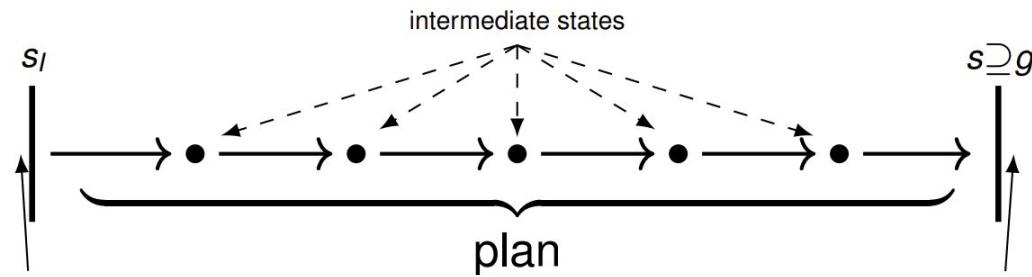
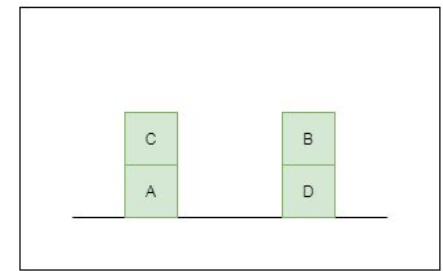
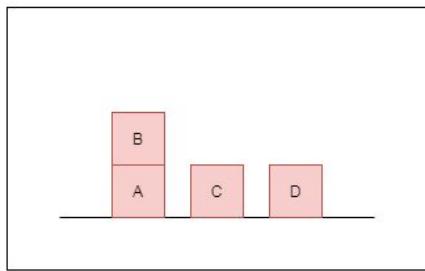
description of desired
world properties

Diagram by
Pascal Bercher

Classical planning: discrete, deterministic, fully observable and single agent.

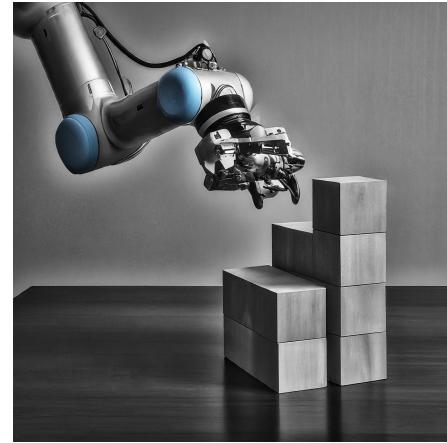
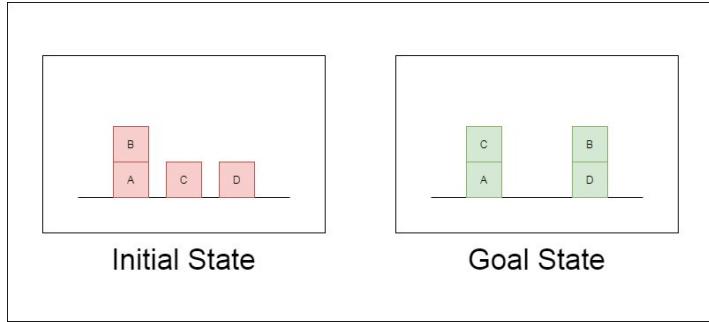
Problem description induces a state transition system, with nodes as states and actions as transitions.

Becomes very complex very quickly: state space is $2^{|V|}$, where V is the number of state features



Pickup B -> Stack B on D -> Pickup C -> Stack C on A

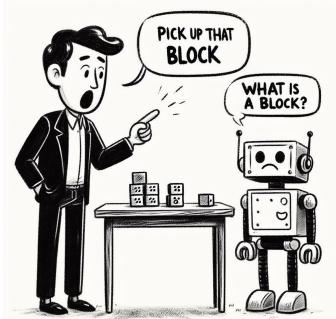




How does it know what a block is ?

How does it know how to stack, pick up and put down?

How does it know what's the best way to pick up?



Outline

Motivation for integrating task and motion planning

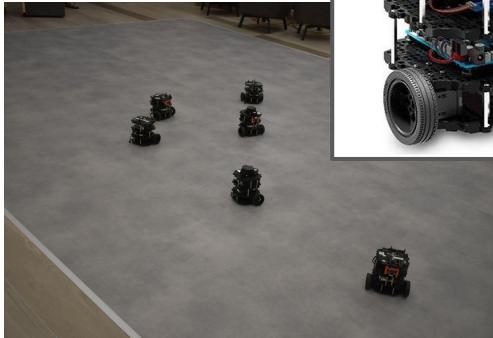
Components of a robotic agent

Task planning and its limitations

Motion planning and its limitations

Integrating task and motion planning

Many Types of Robots

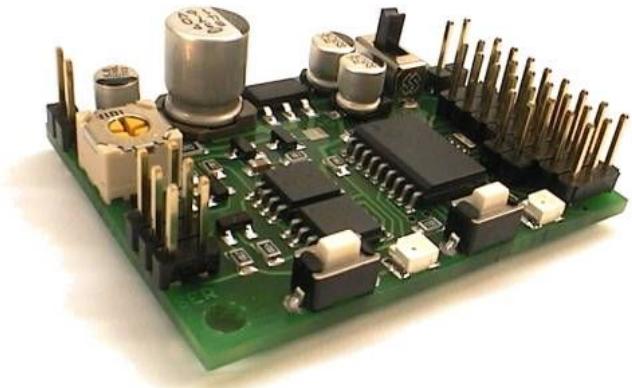


Robot Elements

Sensors



Controllers



Actuators



Robotic Agents



state estimation

Maintain a **belief** $\beta \in \boldsymbol{\beta}$ as a probability distribution over world states S

$$\beta: S \rightarrow [0,1]$$

decision making

Find a **policy**, mapping **belief** β and **objective** into actions (probabilities)

$$\pi: \boldsymbol{\beta} \square A \rightarrow [0,1]$$

motion control

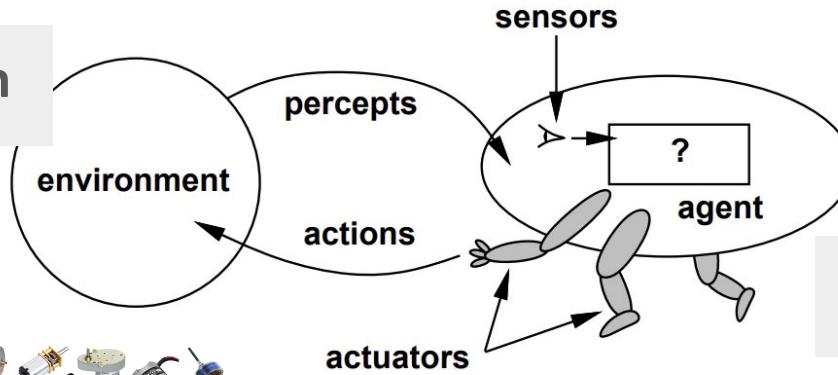
Translate actions into low-level commands (and monitor their execution)

$$u = k_p \cdot e + k_i \cdot \int e dt + k_d \cdot \frac{d}{dt} e$$



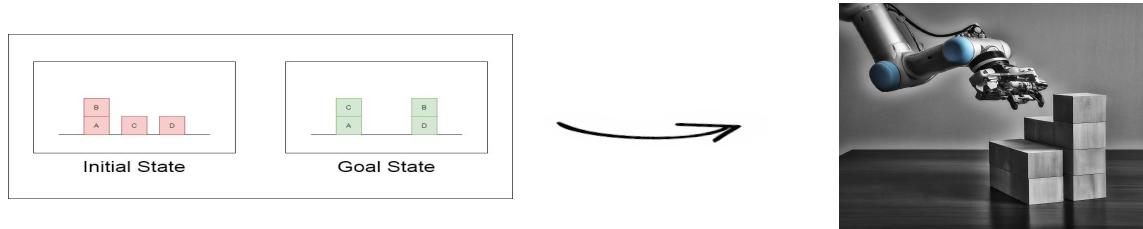
Robotic Agents

state estimation



decision making

motion control



decision making

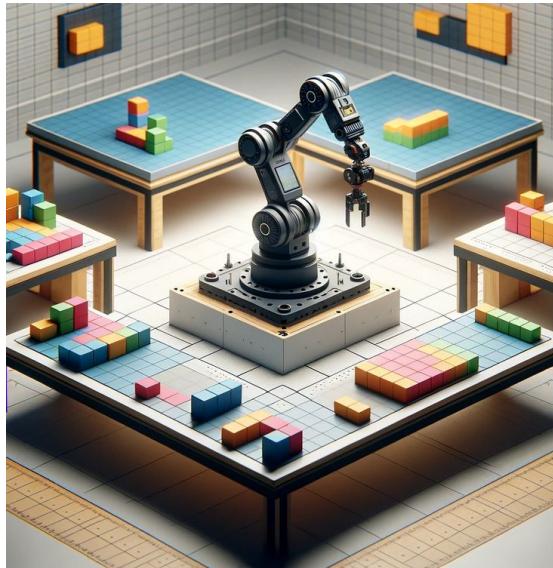
```
(:action stack
  :parameters (?b1 ?b2 - block)
  :precondition (and (holding ?b1) (clear ?b2))
  :effect
    (and (not (holding ?b1))
         (not (clear ?b2))
         (clear ?b1)
         (handempty)
         (on ?b1 ?b2)))
```

state estimation

motion control

Task planning is not enough

Things get messier: N-table Blocks World



Accounting for varying costs, collision avoidance, motion feasibility, multiple (infinite) positions from which an action can be performed, multiple (infinite) ways to reach a position

Outline

Motivation for integrating task and motion planning

Components of a robotic agent

Task planning and its limitations

Motion planning and its limitations

Integrating task and motion planning

Task Planning - Classical Formulation

$$\langle S, s_0, S_G, A, f, c \rangle$$

State space induced by feature set V

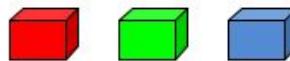
- S is a finite and discrete set of states,
- $s_0 \in S$ is the *known initial state*,
- $S_G \subseteq S$ is the non-empty set of goal states,
- $A(s) \subseteq A$ represents the set of actions in A that are applicable in each state $s \in S$,
- $f(a, s)$ is the *deterministic transition function* where $s' = f(a, s)$ is the state that follows s after doing action $a \in A(s)$, and
- $c(a, s)$ is a *positive cost* for doing action a in the state s .

A solution or *plan* in this model is a sequence of applicable actions a_0, \dots, a_n that generates a state sequence s_0, s_1, \dots, s_{n+1} where s_{n+1} is a goal state. More precisely, the action a_i is applicable

Task Planning - Classical Formulation

Action/model:

Objects (things):



Predicates (true or false statements about the world, "facts"):

(On A, B)

(Clear A)

Actions (how to change the predicates):

(Clear A)

(Clear B)

Stack

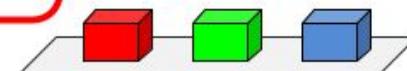
¬(Clear B)

(On A, B)

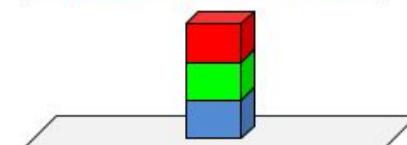
preconditions

effects

Initial: (Clear Red) (Clear Green) (Clear Blue)



Goal: (On Red, Green) (On Green, Blue)



Task Planning: Solution approaches

Common algorithms/techniques:

- Search – forward, regression, bidirectional, symbolic
- Compilation – to SAT, to CP
- Partial Order Planning
- Dynamic programming: Bellman Backup – i.e. for problems with uncertainty and rewards, like MDPs and POMDPs.

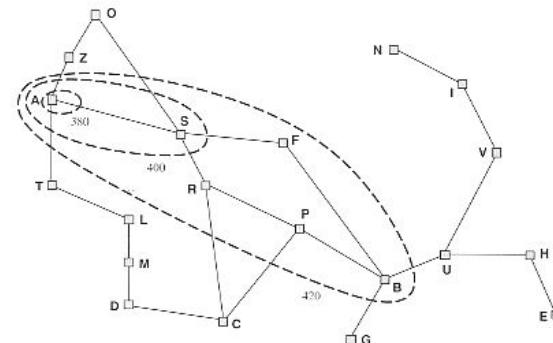
What are some general search algorithms?

Best First Search

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
    node  $\leftarrow$  NODE(STATE=problem.INITIAL)
    frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
    reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
    while not IS-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s]  $\leftarrow$  child
                add child to frontier
    return failure

function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

In our implementation, we check if the node is terminal



Code above from Norvig, P. Russel, and S. Artificial Intelligence. "A modern approach." Prentice Hall Upper Pearl, Judea. Heuristics: intelligent search strategies for computer problem solving. Addison-Wesley Longman Publishing Co., Inc., 1984.

What is not captured by classical task planning ?

Task Planning - Extensions

Accounting for stochastic action outcome
using a state transition probability function

$$\mathcal{P}_{s,s'}^a = \mathcal{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

Accounting for objectives using a reward function

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

Accounting for reward depreciation using a discount factor

$$\gamma \in [0, 1]$$

Accounting for partial observability using
observations Ω and observation function

$$\mathcal{O}_{s,a}^o = \mathcal{P}[O_{t+1} = o | S_t = s, A_t = a]$$

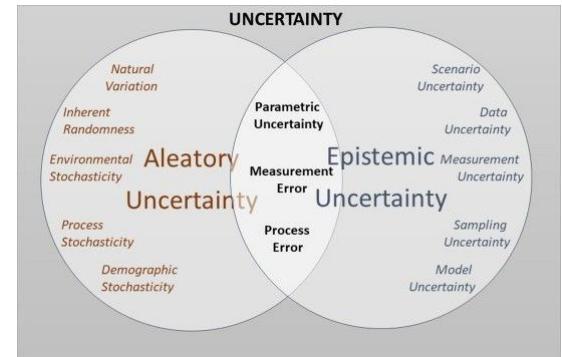
Accounting for action duration, multiple agents, varying costs, etc.

Sequential decision making: what should the agent do to maximize total rewards?

What is not captured by task planning ?

Task Planning vs. the Real World

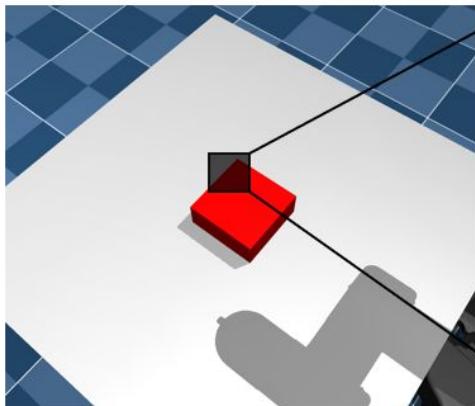
- **Discrete:** plans are finite sequences of actions - while actual actions are continuous
- **Fully modeled actions** (deterministic in classical planning): while actual actions can have unexpected effects that are hard/impossible to model
- **Full knowledge** of world state or of belief (distribution over possible world states) while knowing the actual state /distribution relies on complex sensing and reasoning



https://www.stat.berkeley.edu/users/aldous/157/Papers/Fox_Ulkumen.pdf

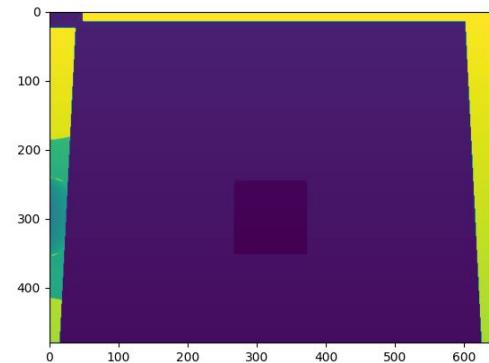
State Estimation

How do we know that block B is on Table 3?



(128, 128, 156)	(128, 62, 53)	(96, 128, 159)	(254, 128, 156)
(122, 124, 200)	(111, 138, 156)	(128, 128, 19)	(252, 128, 156)
(122, 198, 96)	(155, 128, 156)	(242, 128, 156)	(201, 113, 157)
(128, 18, 102)	(255, 181, 156)	(199, 178, 156)	(129, 176, 196)

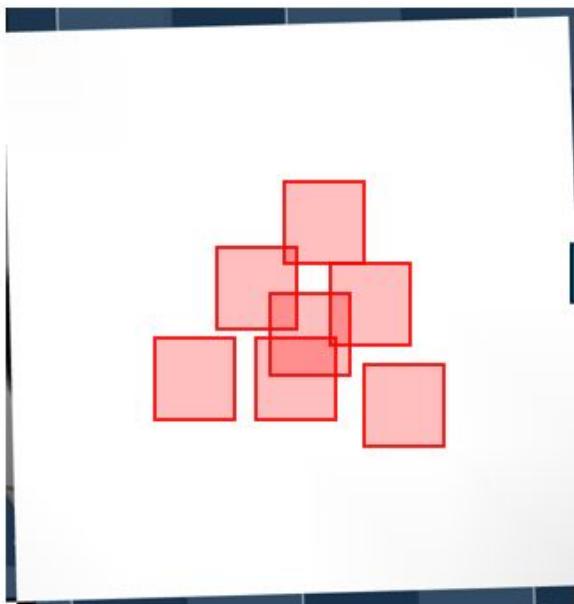
RGB image



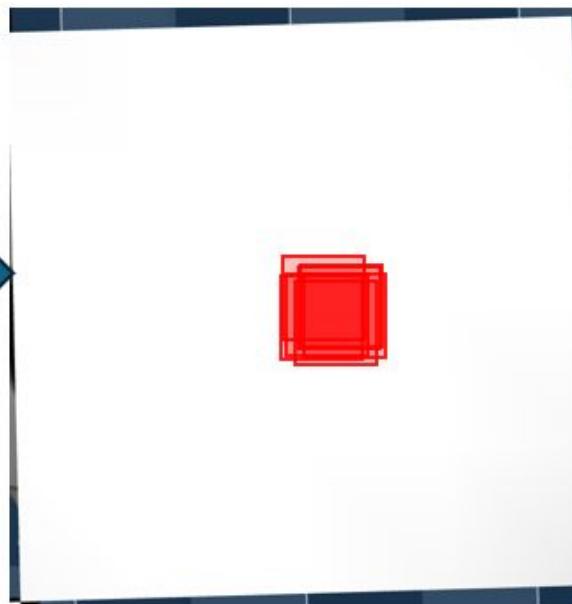
Depth image

Belief Update

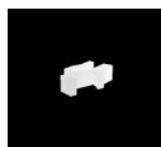
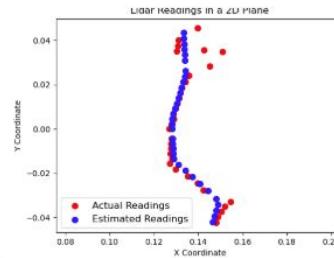
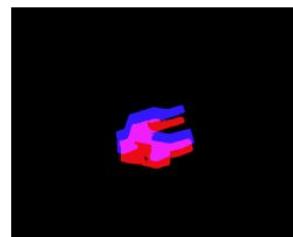
β



β'



Belief Update



(a) P1

(b) P2

(c) P3

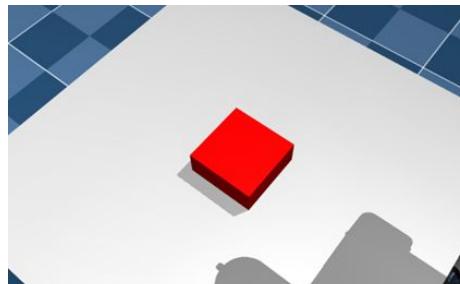
(d) P4

(e) P5

(f) P6

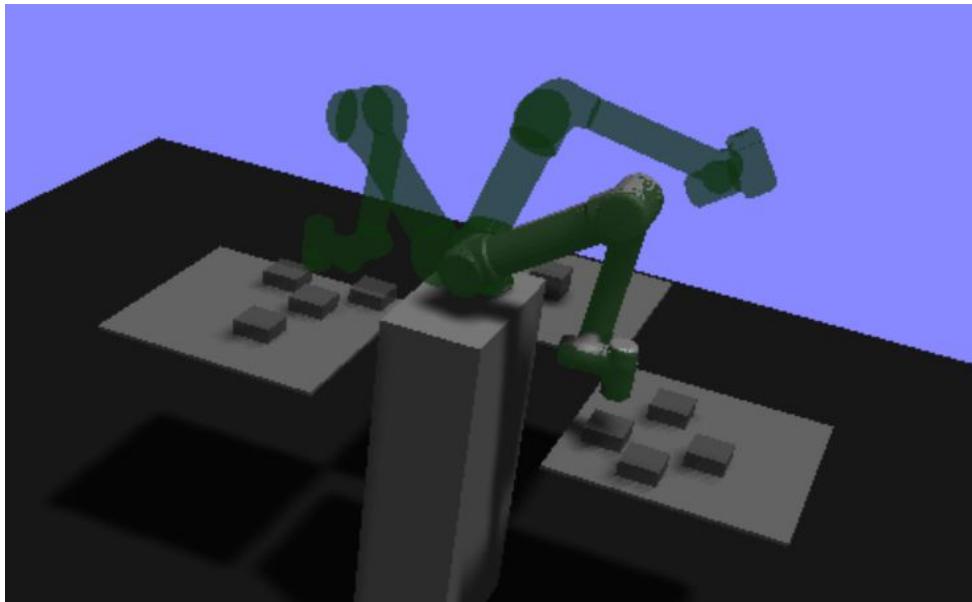
Motion Planning

How to attempt to pick up block B ? How to reach that position ?



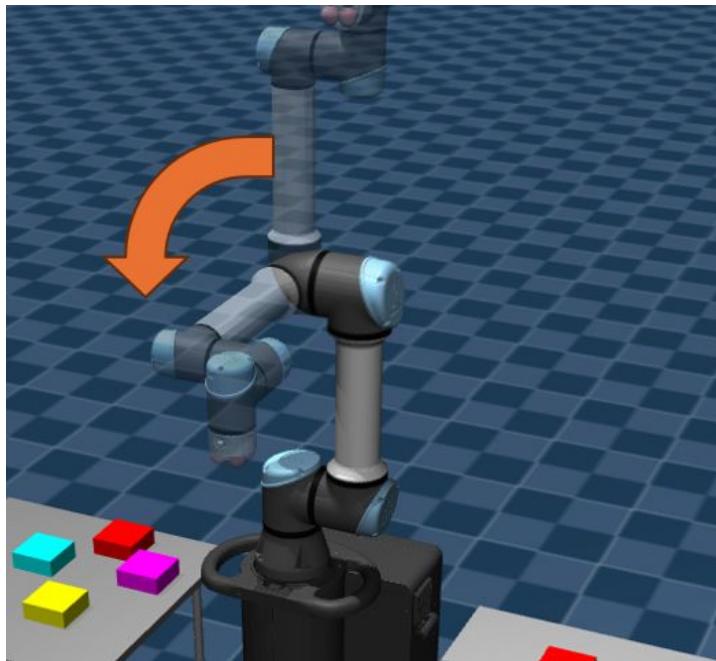
Motion Planning

How to get from one position to another ?



Control and Monitoring

how much force to apply?



move fast vs. accurate the motions should be ?

how to verify that the block was actually picked up (and not dropped)?

Outline

Motivation for integrating task and motion planning

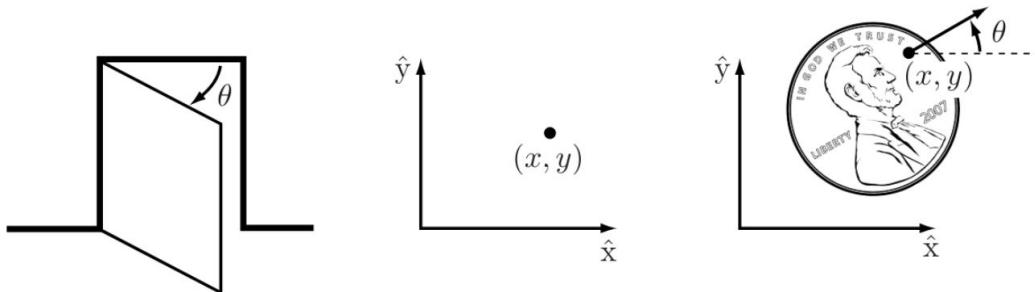
Components of a robotic agent

Task planning and its limitations

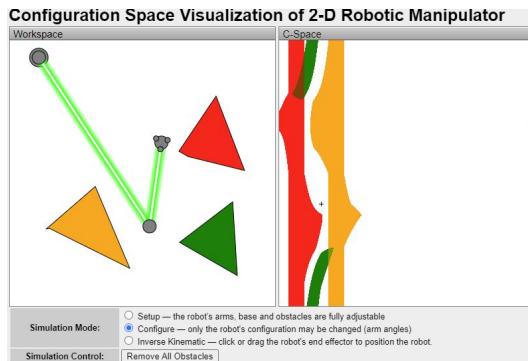
Motion planning and its limitations

Integrating task and motion planning

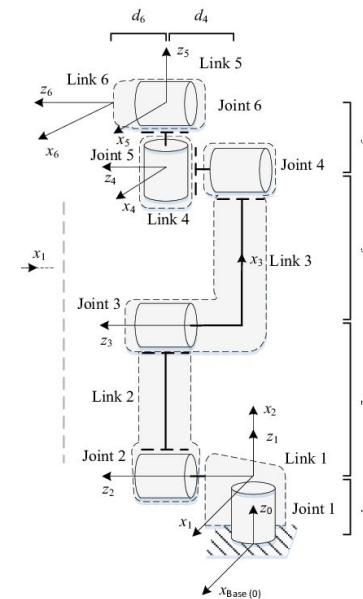
Configuration Space and Degrees of Freedom (DOF)



<https://hades.mech.northwestern.edu/images/7/7f/MR.pdf>

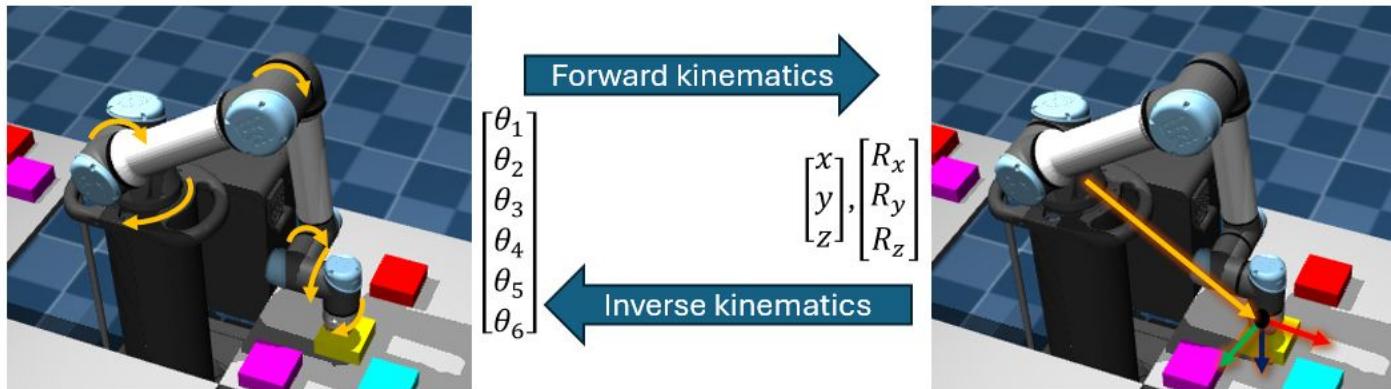


<https://www.cs.unc.edu/~jeffi/c-space/robot.xhtml>



Galin, Rinat, Roman Meshcheryakov, and Anna Samoshina. "Mathematical modelling and simulation of human-robot collaboration." *2020 International Russian Automation Conference (RusAutoCon)*. IEEE, 2020.

Configuration Space

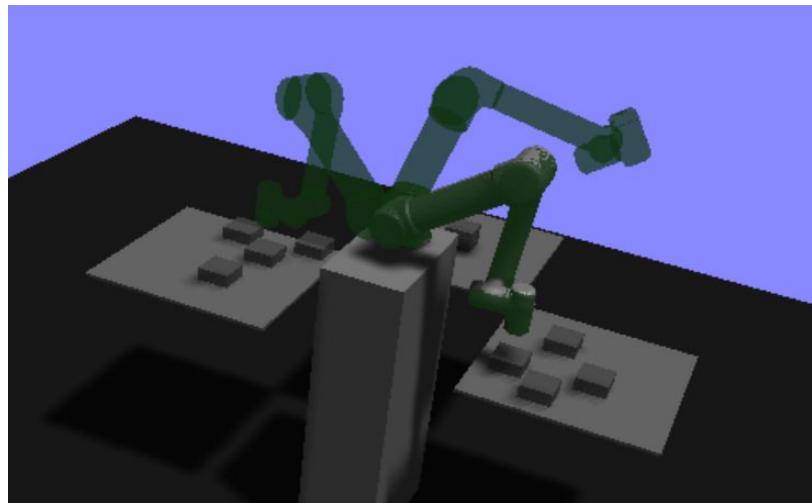


Translating between joint states and end effector pose

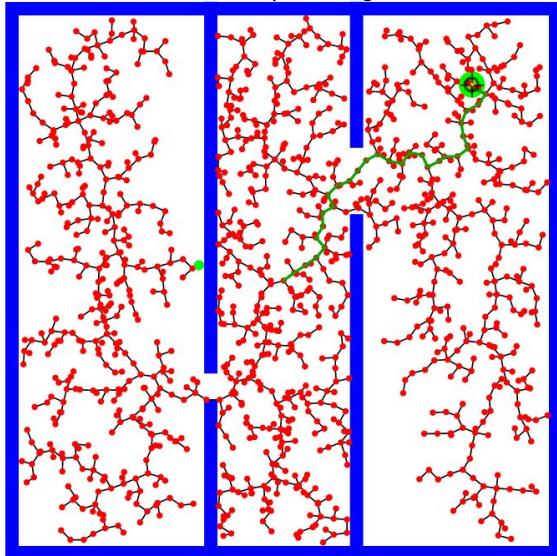
Motion Planning Problem in configuration space (C-Space)

Given configurations $x_i, x_g \in X$, compute a path from x_i to x_g

Path: continuous function $\tau: [0,1] \rightarrow X$ s.t. $\tau(0) = x_i$ and $\tau(1) = x_g$ and $\forall t \in [0,1] \tau(t) \text{ not in } X_{\text{obs}}$

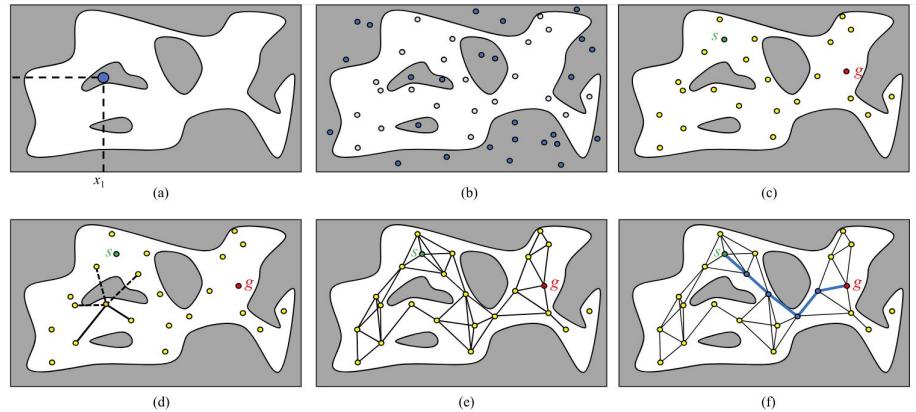


Many approaches (we will see some in the labs)



Rapidly Exploring Random Trees

<https://demonstrations.wolfram.com/RapidlyExploringRandomTreeRRTAndRRT/>



Probabilistic Road Map (PRM)

<http://motion.cs.illinois.edu/RoboticSystems/figures/planning/prm.svg>

What is not captured by motion planning ?

Task planning vs. motion planning

Task planners can reason over very large sets of states by manipulating partial descriptions, while geometric **motion planners** operate on completely detailed specifications of world states.

A task planner could decide that the living room needs to be traversed, disregarding the detailed arrangement of its furniture.

Motion planners deal beautifully with geometry, but not with non-physical aspects of the domain; they can plan how to get to the phone but not decide that a phone call needs to be made.

Explanation as I heard it from Leslie Kaelbling

Outline

Motivation for integrating task and motion planning

Components of a robotic agent

Task planning and its limitations

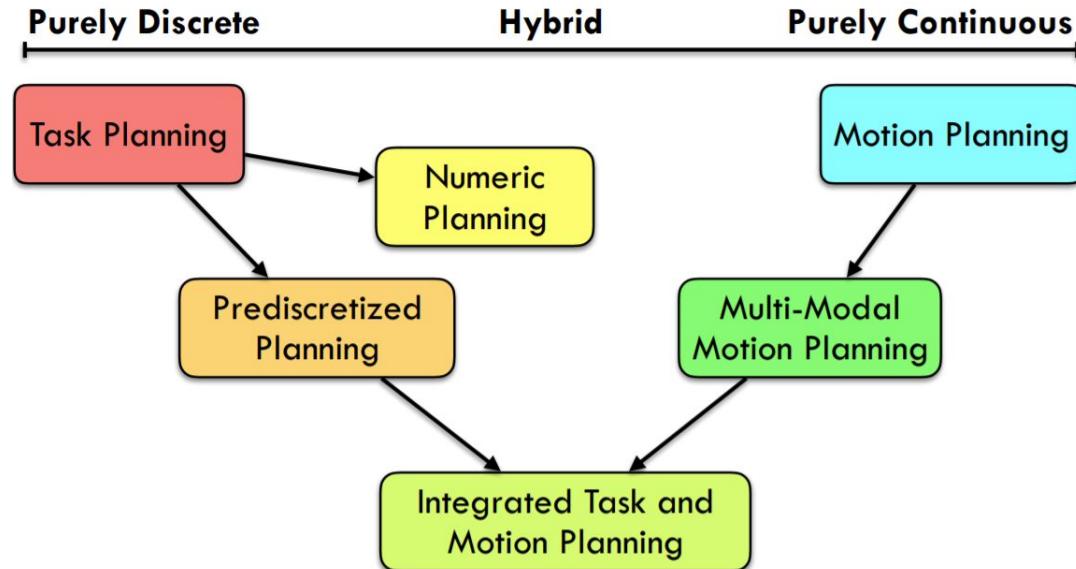
Motion planning and its limitations

Integrating task and motion planning

Task and Motion Planning: Challenges

- Inherits challenges of both motion & task planning
 - high-dimensional, continuous state-spaces
 - long horizons
- Geometric and continuous constraints limit high-level strategies
 - kinematics, reachability, joint limits, collisions, grasp, visibility, stability, stiffness, torque limits,

Approaches to Task and Motion Planning



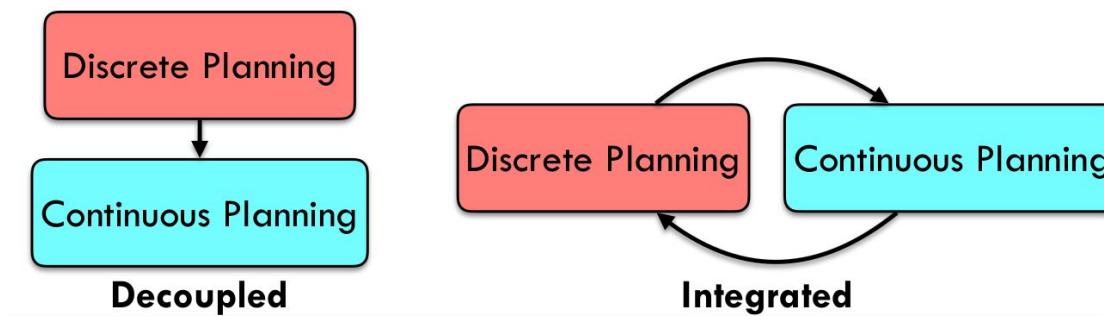
Decoupled vs. Integrated Task and Motion Planning

Decoupled:

- discrete (task) planning then continuous (motion) planning
- requires a strong **downward refinement assumption**: every correct discrete plan can be refined into a correct continuous plan (hierarchical planning)

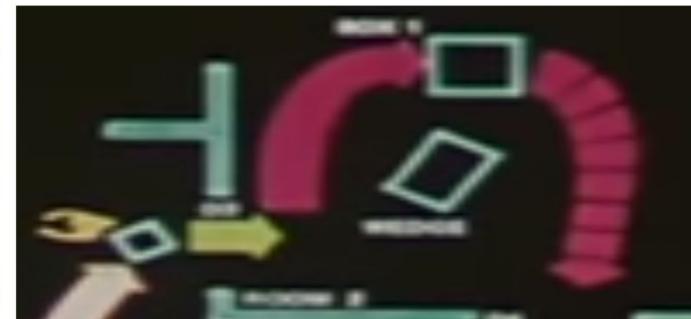
Integrated:

- simultaneous discrete & continuous planning



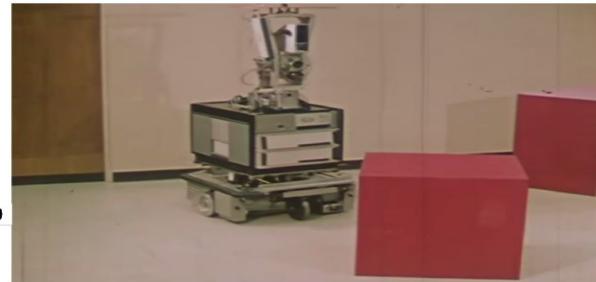
Shakey the Robot (1969): Task planning then motion planning

- First autonomous mobile manipulator (via pushing)
- Visibility graph, A* search, and STRIPS (which is a grounded form of PDDL) [Fikes 1971, Nilsson 1984]



```
type(robot robot) type(ol object)
name(robot shakey) name(ol box1)
at(robot 4.1 7.2) at(ol 3.1 5.2)
theta(robot 90.1) inroom(ol r1)
shape(ol wedge)
radius(ol 3.1)

GOTHRU(d,r1,r2)
Precondition INROOM(ROBOT,r1) ∧ CONNECTS(d,r1,r2)
Delete List INROOM(ROBOT,$)
Add List INROOM(ROBOT,r2)
```



Can we suggest better ways ?